

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import statsmodels.api as sm
```

```
df = pd.read_csv("marketing_and_sales_data_evaluate_lr.csv")

df.head()
```

	TV	Radio	Social_Media	Sales
0	16.0	6.566231	2.907983	54.732757
1	13.0	9.237765	2.409567	46.677897
2	41.0	15.886446	2.913410	150.177829
3	83.0	30.020028	6.922304	298.246340
4	15.0	8.437408	1.405998	56.594181

```
print(df.shape)

print(df.info())

print(df.isnull().sum())
```

```
(4572, 4)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4572 entries, 0 to 4571
Data columns (total 4 columns):
#   Column          Non-Null Count  Dtype
---  -
0    TV              4562 non-null   float64
1    Radio           4568 non-null   float64
2    Social_Media    4566 non-null   float64
3    Sales           4566 non-null   float64
dtypes: float64(4)
memory usage: 143.0 KB
None
TV              10
Radio           4
Social_Media    6
Sales           6
dtype: int64
```

```
df = df.dropna()
```

```
print(df.isnull().sum())
```

```
TV              0
Radio           0
Social_Media    0
Sales           0
dtype: int64
```


```
correlation_matrix = df.corr(numeric_only=True)

print(correlation_matrix)

sns.heatmap(correlation_matrix,
            annot=True,
            cmap="coolwarm")


plt.title("Correlation Matrix")
plt.show()
```

	TV	Radio	Social_Media	Sales
TV	1.000000	0.869158	0.527687	0.999497
Radio	0.869158	1.000000	0.606338	0.868638
Social_Media	0.527687	0.606338	1.000000	0.527446
Sales	0.999497	0.868638	0.527446	1.000000



```
sns.pairplot(df)
```

```
plt.show()
```



```
X = df["TV"]
y = df["Sales"]
X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
print(model.summary())
```

```
=====
                        OLS Regression Results
=====
Dep. Variable:          Sales      R-squared:                0.999
Model:                  OLS       Adj. R-squared:           0.999
Method:                 Least Squares   F-statistic:           4.517e+06
Date:                   Thu, 11 Jun 2026   Prob (F-statistic):       0.00
Time:                   07:17:31         Log-Likelihood:         -11366.
No. Observations:       4546            AIC:                   2.274e+04
Df Residuals:           4544            BIC:                   2.275e+04
Df Model:                1
Covariance Type:        nonrobust
=====
coef    std err          t      P>|t|    [0.025    0.975]
-----
-----
```

```

const      -0.1325    0.101    -1.317    0.188    -0.330    0.065
TV          3.5615    0.002    2125.272    0.000    3.558    3.565
=====
Omnibus:                0.052    Durbin-Watson:                1.998
Prob(Omnibus):           0.974    Jarque-Bera (JB):           0.031
Skew:                   -0.001    Prob(JB):                 0.985
Kurtosis:                3.012    Cond. No.                  138.
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
predictions = model.predict(X)
```


```
residuals = y - predictions
```

```

sns.regplot(
    x=df["TV"],
    y=df["Sales"]
)

plt.title("TV vs Sales")
plt.show()

```



```

plt.scatter(predictions, residuals)


plt.axhline(y=0)

plt.xlabel("Predicted Sales")
plt.ylabel("Residuals")

plt.title("Residual Plot")

plt.show()

```



```
sm.qqplot(residuals, line="45")
```

```
plt.show()
```